

# Introduction to RAG and Agentic RAG with LlamalIndex and LangChain

AI Engineering: Agents, Vibe Coding and Full-Stack AI Course

John Alexander




Slides and code repo: <https://github.com/JRAlexander/Intro-to-RAG-Agentic-RAG-2602>

# Why RAG?

# Large Language Model challenges

- Trained with point-in-time data to ensure effectiveness
- Concerns about costs and security
- Hallucinations – confidently wrong answers


# Integrating domain knowledge



**RAG 101**

# What is Retrieval Augmented Generation (RAG)

- a pattern that works with pretrained Large Language Models (LLM) and your own data to generate responses.
- Use tools and components to augment the prompting!




# DEMO

Simple RAG

# Why do we need to split documents?

- 1 LLMs have limited context windows (4K – 128K)
- 2 When an LLM receives too much information, it can get easily distracted by irrelevant details.
- 3 The more tokens you send, the higher the cost, the slower the response.



Source: Pamela Fox, <https://aka.ms/pythonai/slides/rag>

# Optimal size of document chunk

## How big should chunks be?

| # of tokens per chunk | Recall@50 |
|-----------------------|-----------|
| 512                   | 42.4      |
| 1024                  | 37.5      |
| 4096                  | 36.4      |
| 8191                  | 34.9      |

Source: [aka.ms/ragrelevance](https://aka.ms/ragrelevance)

A **token** is the unit of measurement for an LLM's input/output. ~1 token/word for English, higher ratios for other languages.

More on token ratios: [aka.ms/genai-cjk](https://aka.ms/genai-cjk)

## Where to split chunks?

| Chunk boundary strategy      | Recall@50 |
|------------------------------|-----------|
| Break at token boundary      | 40.9      |
| Preserve sentence boundaries | 42.4      |
| 10% overlapping chunks       | 43.1      |
| 25% overlapping chunks       | 43.9      |

Source: [aka.ms/ragrelevance](https://aka.ms/ragrelevance)

A chunking algorithm should also consider tables, and avoid splitting tables when possible.

Source: Pamela Fox, <https://aka.ms/pythonai/slides/rag>

# Tokens – the coin of the realm

- Models process text using **tokens**

Many words map to one token, but some don't: indivisible.

Unicode characters like emojis may be split into many tokens containing the underlying bytes: ♡♦♦♦♦

Sequences of characters commonly found next to each other may be grouped together: 1234567890



```
[7085, 2456, 3975, 284, 530, 11241, 11, 475, 617, 836, 470, 25, 773, 452, 12843, 13, 198, 198, 3118, 291, 1098, 3435, 588, 795, 13210, 271, 743, 307, 6626, 656, 867, 16326, 7268, 262, 10238, 9881, 25, 12520, 97, 248, 8582, 237, 122, 198, 198, 44015, 3007, 286, 3435, 8811, 1043, 1306, 284, 1123, 584, 743, 307, 32824, 1978, 25, 17031, 2231, 30924, 3829]
```


Each token counts towards the model's maximum limit, so it's important to consider the token count when designing your prompt (and tool calls and context).

# DEMO

Chunking




# Introduction to Vector Databases

---



# What are Vector Embeddings?

- a vector (list) of floating-point numbers
- measure the relatedness of text strings



- The distance between two vectors measures their relatedness.
- Small distances suggest high relatedness and large distances suggest low relatedness.

a quarterback  
throws a football

# Embedding and Text Similarity

feline friends say

meow


canine companions say

woof

bovine buddies say

moo


# What is a Vector Database?



# DEMO

Embedding and Text Similarity

# RAG data source types



## Database rows (Structured data)

You need a way to **vectorize** target columns with an **embedding model**.

You need a way to **search** the vectorized rows.



## Documents (Unstructured data)


PDFs, docx, pptx, md, html, images

You need an ingestion process for **extracting, splitting, vectorizing, and storing** document chunks.

You need a way to **search** the vectorized **chunks**.

# RAG document ingestion

For long/unstructured documents, we need an ingestion flow such as this one:



## Extract text from PDF

Other options for this step:  
Azure Document Intelligence,  
LlamaParse,  
LangChain document loaders,  
OCR services, Unstructured, etc.

## Split data into chunks

Split text based on sentence boundaries and token lengths.  
You could also use "semantic" splitters and your own custom splitters.

## Vectorize chunks

Compute embeddings using embedding model of your choosing.


## Store chunks

This is where you'd typically use a search service like Azure AI Search or a database like PostgreSQL.

# Orchestration Tools

# Introduction to Llamalndex

- Designed for efficient information retrieval and data analysis
- Optimized to handle large datasets and provide fast, accurate data access
- Built to scale with increasing data sizes and complexity, making it suitable for various applications
- Supports many different data types and retrieval methods




# LlamalIndex: Real-World Applications

Used in applications such as:

- semantic search
- contextual data retrieval, and real-time data analysis.

# Introduction to LangChain

- Powerful framework designed for creating intelligent pipelines for text generation and processing.
- Leverages the capabilities of large language models
- Chain = pipeline
- Large number of tools and frameworks integrations
- Modular design allows easy customization and integration of different components.



# LangChain: Real-World Applications

Used in various applications such as:

- automated report generation
- intelligent document summarization
- real-time sentiment analysis


# Introduction to LangGraph

- Library for building stateful, multi-actor applications with LLMs, used to create agent and multi-agent workflows
- Provides fine-grained control over both the flow and state of your agent applications
- Agentic Features
  - Human-in-the-loop
  - Parallelization
  - Subgraphs
  - Reflection

# When to use LangGraph

- You need fine-grained, low-level control over agent orchestration.
- You need durable execution for long-running, stateful agents.
- You're building complex workflows that combine deterministic and agentic steps.
- You need production-ready infrastructure for agent deployment.

# Agentic RAG – LangChain / LangGraph Architecture



# Why plain RAG breaks

- Retrieves irrelevant chunks (keyword bias, poor query)
- Needs multiple hops ( $A \rightarrow B \rightarrow C$ ), but retrieval is one-shot
- Conflicts across sources are not reconciled
- Answer quality depends on luck: one query, one retrieval, one pass

# Combining LlamalIndex and LangChain

Agentic RAG Demo

# Workflow

- Step 1: Define the data retrieval task with Llamalndex.
- Step 2: Process and analyze the retrieved data.
- Step 3: Generate answer using LangChain.

# Questions?